

Understanding the MapReduce Programming Model

INTRODUCING MAPREDUCE



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Summary

Understand the need for Distributed Computing

Understand the role of MapReduce in a distributed computing setup

Spot applications of MapReduce

Know the typical flow of a MapReduce task

What Do These Products Have in Common?

**Facebook
Social
Network**

**Google
Search**

**LinkedIn
Member
Analytics**

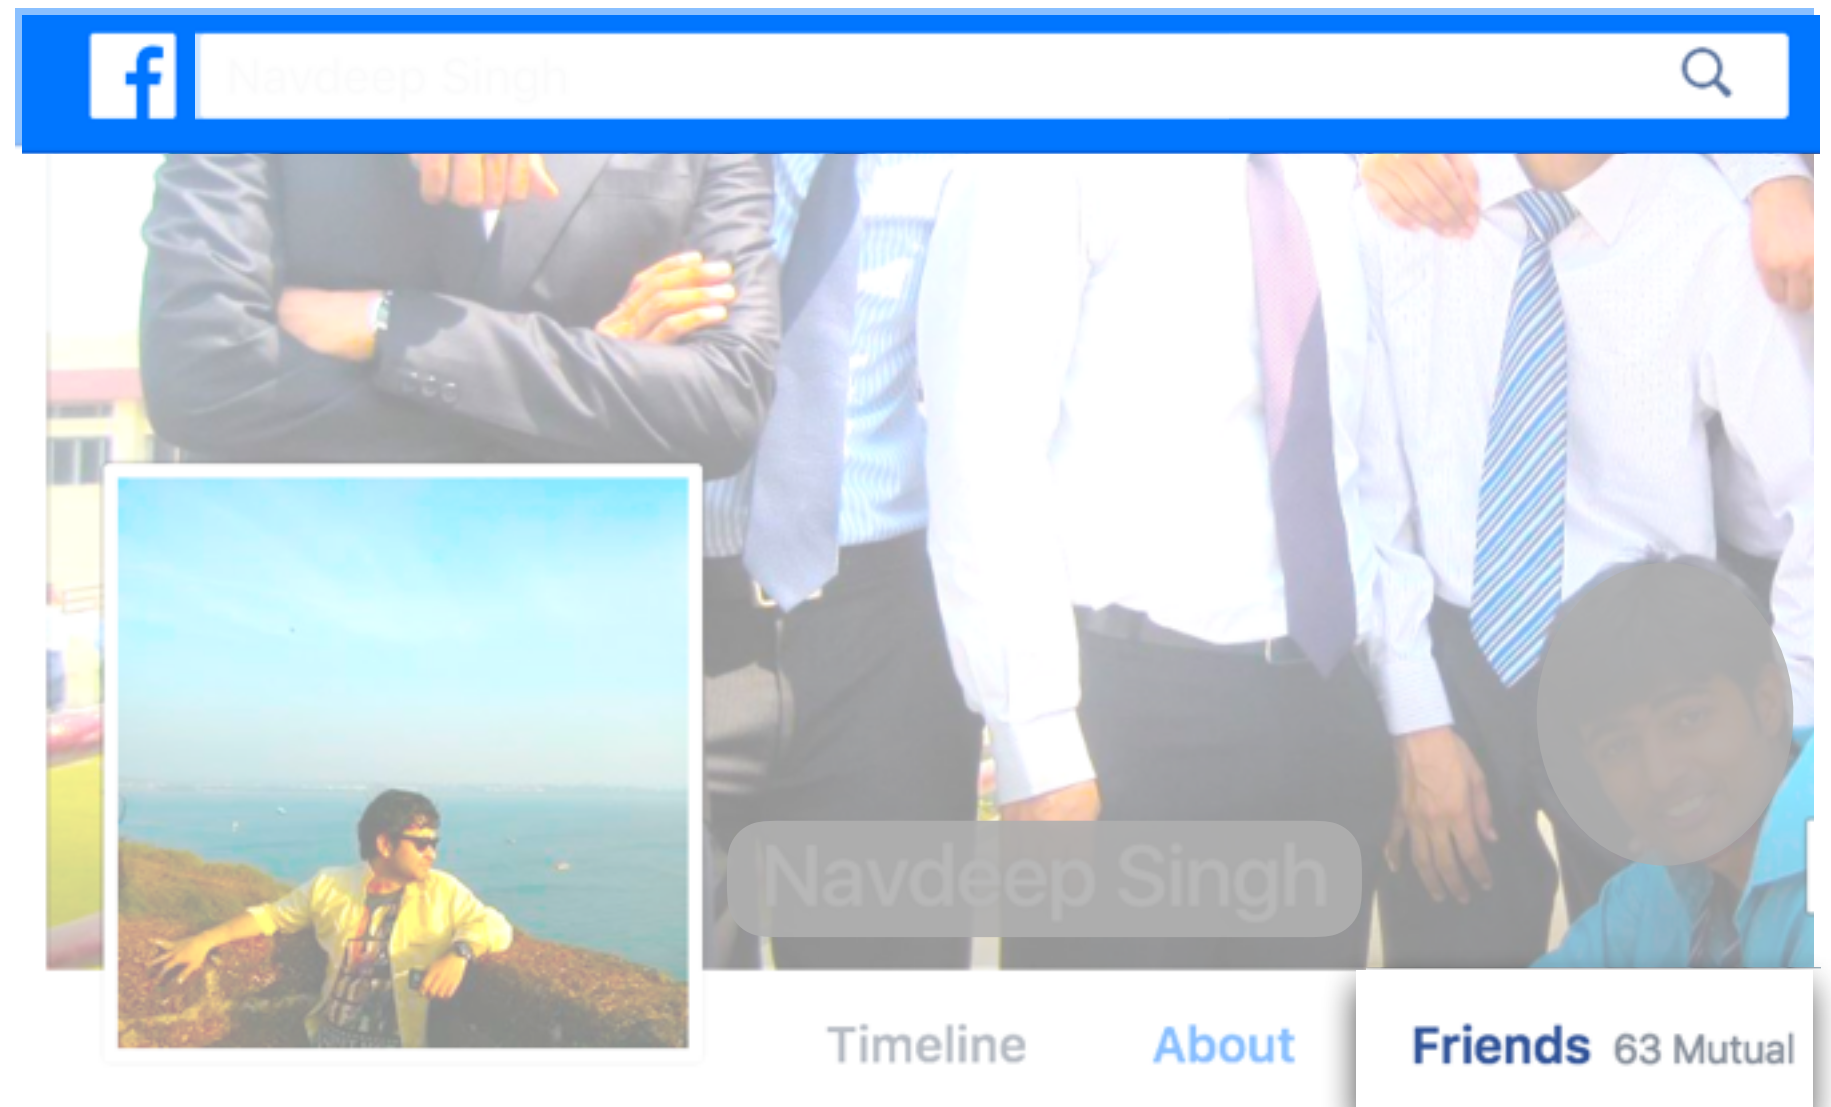
Facebook Social Network

More than a billion users log on to Facebook every month..



Facebook Social Network

Yet, Facebook takes milliseconds
to return the number of mutual
friends for every pair of users



Facebook
Social
Network

Google
Search

LinkedIn
Member
Analytics

Google
Search

The Internet has more than 40 billion webpages




The size of the World Wide Web:
Estimated size of Google's index



Google
Search

Yet, Google can
return a search result
in less than a second



**Facebook
Social
Network**

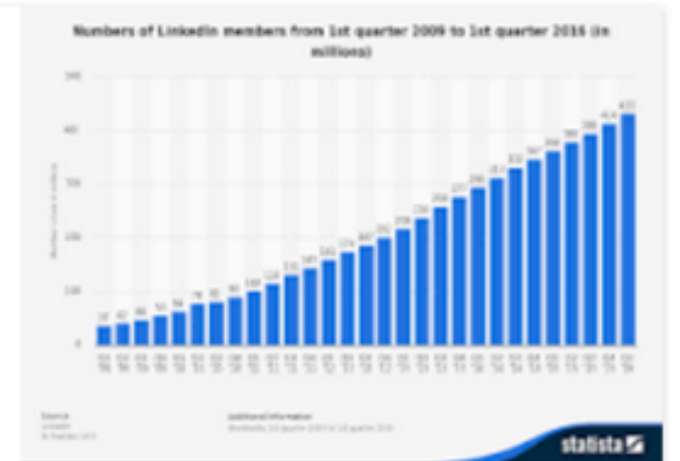
**Google
Search**

**LinkedIn
Member
Analytics**

LinkedIn Member Analytics

LinkedIn has nearly half a billion members..

This timeline displays **member numbers** of social network **LinkedIn** from the first quarter of 2009 to the second quarter of 2016, in millions. During the most recently reported quarter, **LinkedIn** had **450 million members**, up from 414 million **members** in the preceding quarter.



Numbers of LinkedIn members 2009-2016 - Statista

www.statista.com/statistics/274050/quarterly-numbers-of-linkedin-members/

LinkedIn Member Analytics

Yet, LinkedIn can display complicated metrics for each and every member on demand..

7 people viewed your profile in the past 3 days

▲8% profile rank in the past 30 days

What's special about this?

Facebook
Social
Network

Google
Search

LinkedIn
Member
Analytics

Huge data set



One clear insight


7 people viewed your profile in the past 3 days

▲8% profile rank in the past 30 days

The power of indexes!

Billions of rows of raw data

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Swetha



7

people viewed your profile in the past 3 days

▲8%

profile rank in the past 30 days

Building an Index

Billions of rows of
raw data

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Swetha

**Summarize
and Sort**

Index

Member	# Profile Views
Janani	50
Swetha	15
Vitthal	22
Shreya	23
Jitu	32
Pradeep	10

Huge data = Massive Indexes

**An Index is a Map
with a really fast
way to look it up**

Member	# Profile Views
Janani	50
Swetha	15
Vitthal	22
Shreya	23
Jitu	32
Pradeep	10

Massive Indexes

Building an index is

- hugely data intensive
- repetitive

Member	# Profile Views
Janani	50
Swetha	15
Vitthal	22
Shreya	23
Jitu	32
Pradeep	10

Big data systems are all
about building these indexes

What are the requirements
of such a system?

System Requirements

Raw data

**Summarize
and Sort**



Index

System Requirements

Raw data

Summarize
and Sort



Index

**Store massive
amounts of data**



System Requirements

Raw data

**Summarize
and Sort**

Index

Store massive
amounts of data

Process it in a
timely manner

System Requirements

Store

Process


?

Store massive
amounts of data

Process it in a
timely manner

Store
Process
?

Growing size of data



**Store
Process
?**

The infrastructure needs
to keep up with the
growing size of data

System Requirements

Store

Store massive
amounts of data

Process

Process it in a
timely manner

Scale

Scale easily as
data grows

Store

Process

Scale

Distributed Computing Frameworks

like Hadoop were developed for
exactly this

Two Ways to Build a System



Monolithic

Distributed


Two Ways to Build a Team




A star who dribbles and shoots



A team of good players who
know how to pass



A supercomputer



**A cluster of decent machines
that know how to parallelize**

Two Ways to Build a System



Monolithic

Distributed

Monolithic



One star player

Monolithic



**A single
powerful server**

2x Expense

> 2x

Performance

Two Ways to Build a System



Monolithic


Distributed

Distributed




A team of good players who
know how to pass

Distributed




Distributed System



**Many small and
cheap computers
come together...**

**...to act as a
single entity**


Distributed System



Cluster

**Such a system
can scale linearly**

Distributed System



2x Nodes

2x Storage

~ 2x Speed

Store

Process


Scale

Distributed Computing makes for a
lot of complexity

Complexity of Storage and Processing

- **Partitioning**
- **Fault tolerance and recovery**
- **Parallel processing**

Complexity of Storage and Processing



Complexity of coordination across nodes

Complexity of Storage and Processing

Abstraction

**A distributed
computing
framework
abstracts away
these issues**

Complexity of Storage and Processing

The programmer just needs
to *specify what processing
needs to be performed*

...in parallel

Complexity of Storage and Processing

The programmer just needs
to *specify what processing
needs to be performed*

...in parallel

MapReduce

MapReduce

The Power of MapReduce

A solid orange square.

Abstraction

A solid maroon square.

Parallelization

Abstraction

**MapReduce abstracts
away the complexity
into just two operations**

Abstraction

The programmer just
defines 2 functions

map
reduce

Hadoop does the rest -
behind the scenes

Parallelization



map


An operation performed
in parallel, on small
portions of the dataset

Parallelization



map

One Record



Key-Value Output

Parallelization



reduce

An operation to
combine the results of
the map step

Parallelization




reduce

Map
Output

Map
Output

Map
Output

Map
Output



Final Output

MapReduce forces the programmer
to **Think Parallel**

Think Parallel



Filtering
Counting
Ranking
Min/Max/Avg

Whatever the task, break it down into 2 steps

- A step that can be performed in parallel
- A step to combine the intermediate results

Think Parallel

map A step that can be performed in parallel

reduce A step to combine the intermediate results

Think Parallel

**Breaking down any task into
these two steps is almost an art**

**This course will teach you this art -
with lots of opportunities to practice it**

Summary

Distributed Computing scales to deal with data as it grows

MapReduce is an abstraction to express processing tasks in distributed computing

MapReduce helps parallelize expensive tasks like building indexes

LinkedIn displays some interesting metrics when a member visits their profile page

7 people viewed your profile in the past 3 days

7 people viewed your profile in the past 3 days


This is powered by an index

Member	# Profile Views
Janani	50
Swetha	15
Vitthal	22
Shreya	23
Jitu	32
Pradeep	10

Building a User-ViewCount Map

Data is originally captured in this form

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep



User-ViewCount Map

Member	# Profile Views
Janani	50
Swetha	15
Vitthal	22
Shreya	23
Jitu	32
Pradeep	10

MapReduce Flow

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep

**The raw data is really large
(potentially in PetaBytes)**

**It's distributed across many
machines in a cluster**

Each machine holds a **partition of
data**

MapReduce Flow

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani



View ID	From Member	To Member
3	Shreya	Pradeep
4	Jitu	Vitthal



View ID	From Member	To Member
5	Shreya	Janani
6	Jitu	Pradeep






Each partition is given to a different process i.e. to **mappers**

MapReduce Flow

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani

View ID	From Member	To Member
3	Shreya	Pradeep
4	Jitu	Vitthal


View ID	From Member	To Member
5	Shreya	Janani
6	Jitu	Pradeep



**Each mapper
works in parallel**


Map Flow

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani




**Within each mapper, the rows
are processed serially**

Map Flow



Each row emits a {key, value} pair


Map Flow



This operation can occur
independently on each mapper


Map Flow

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani




{Jitu, 1}
{Janani, 1}

View ID	From Member	To Member
3	Shreya	Pradeep
4	Jitu	Vitthal



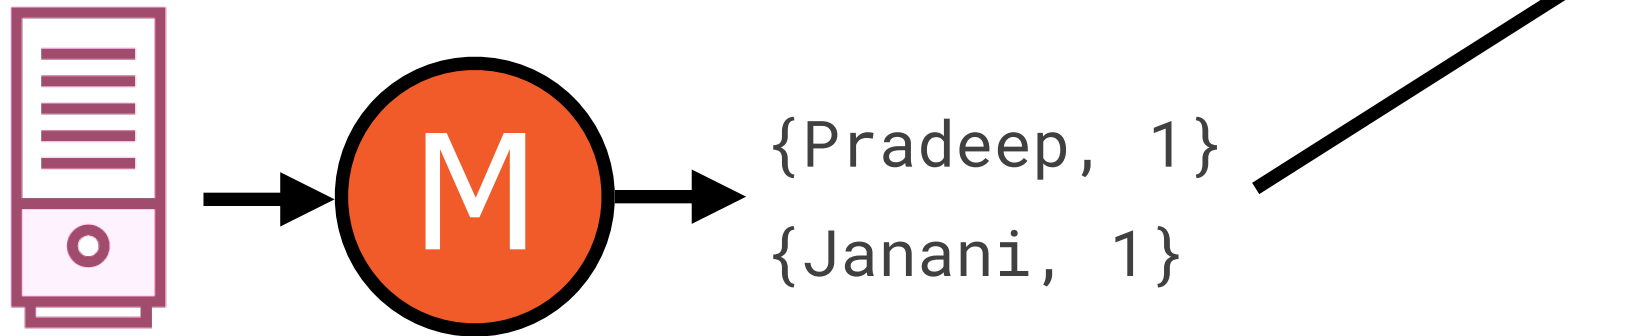

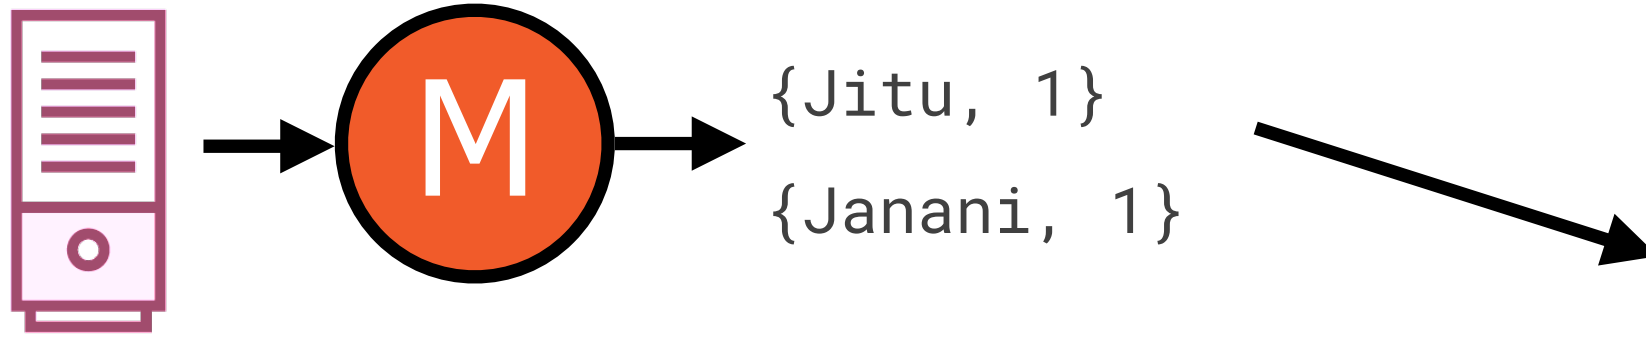
{Pradeep, 1}
{Vitthal, 1}

View ID	From Member	To Member
5	Shreya	Janani
6	Jitu	Pradeep



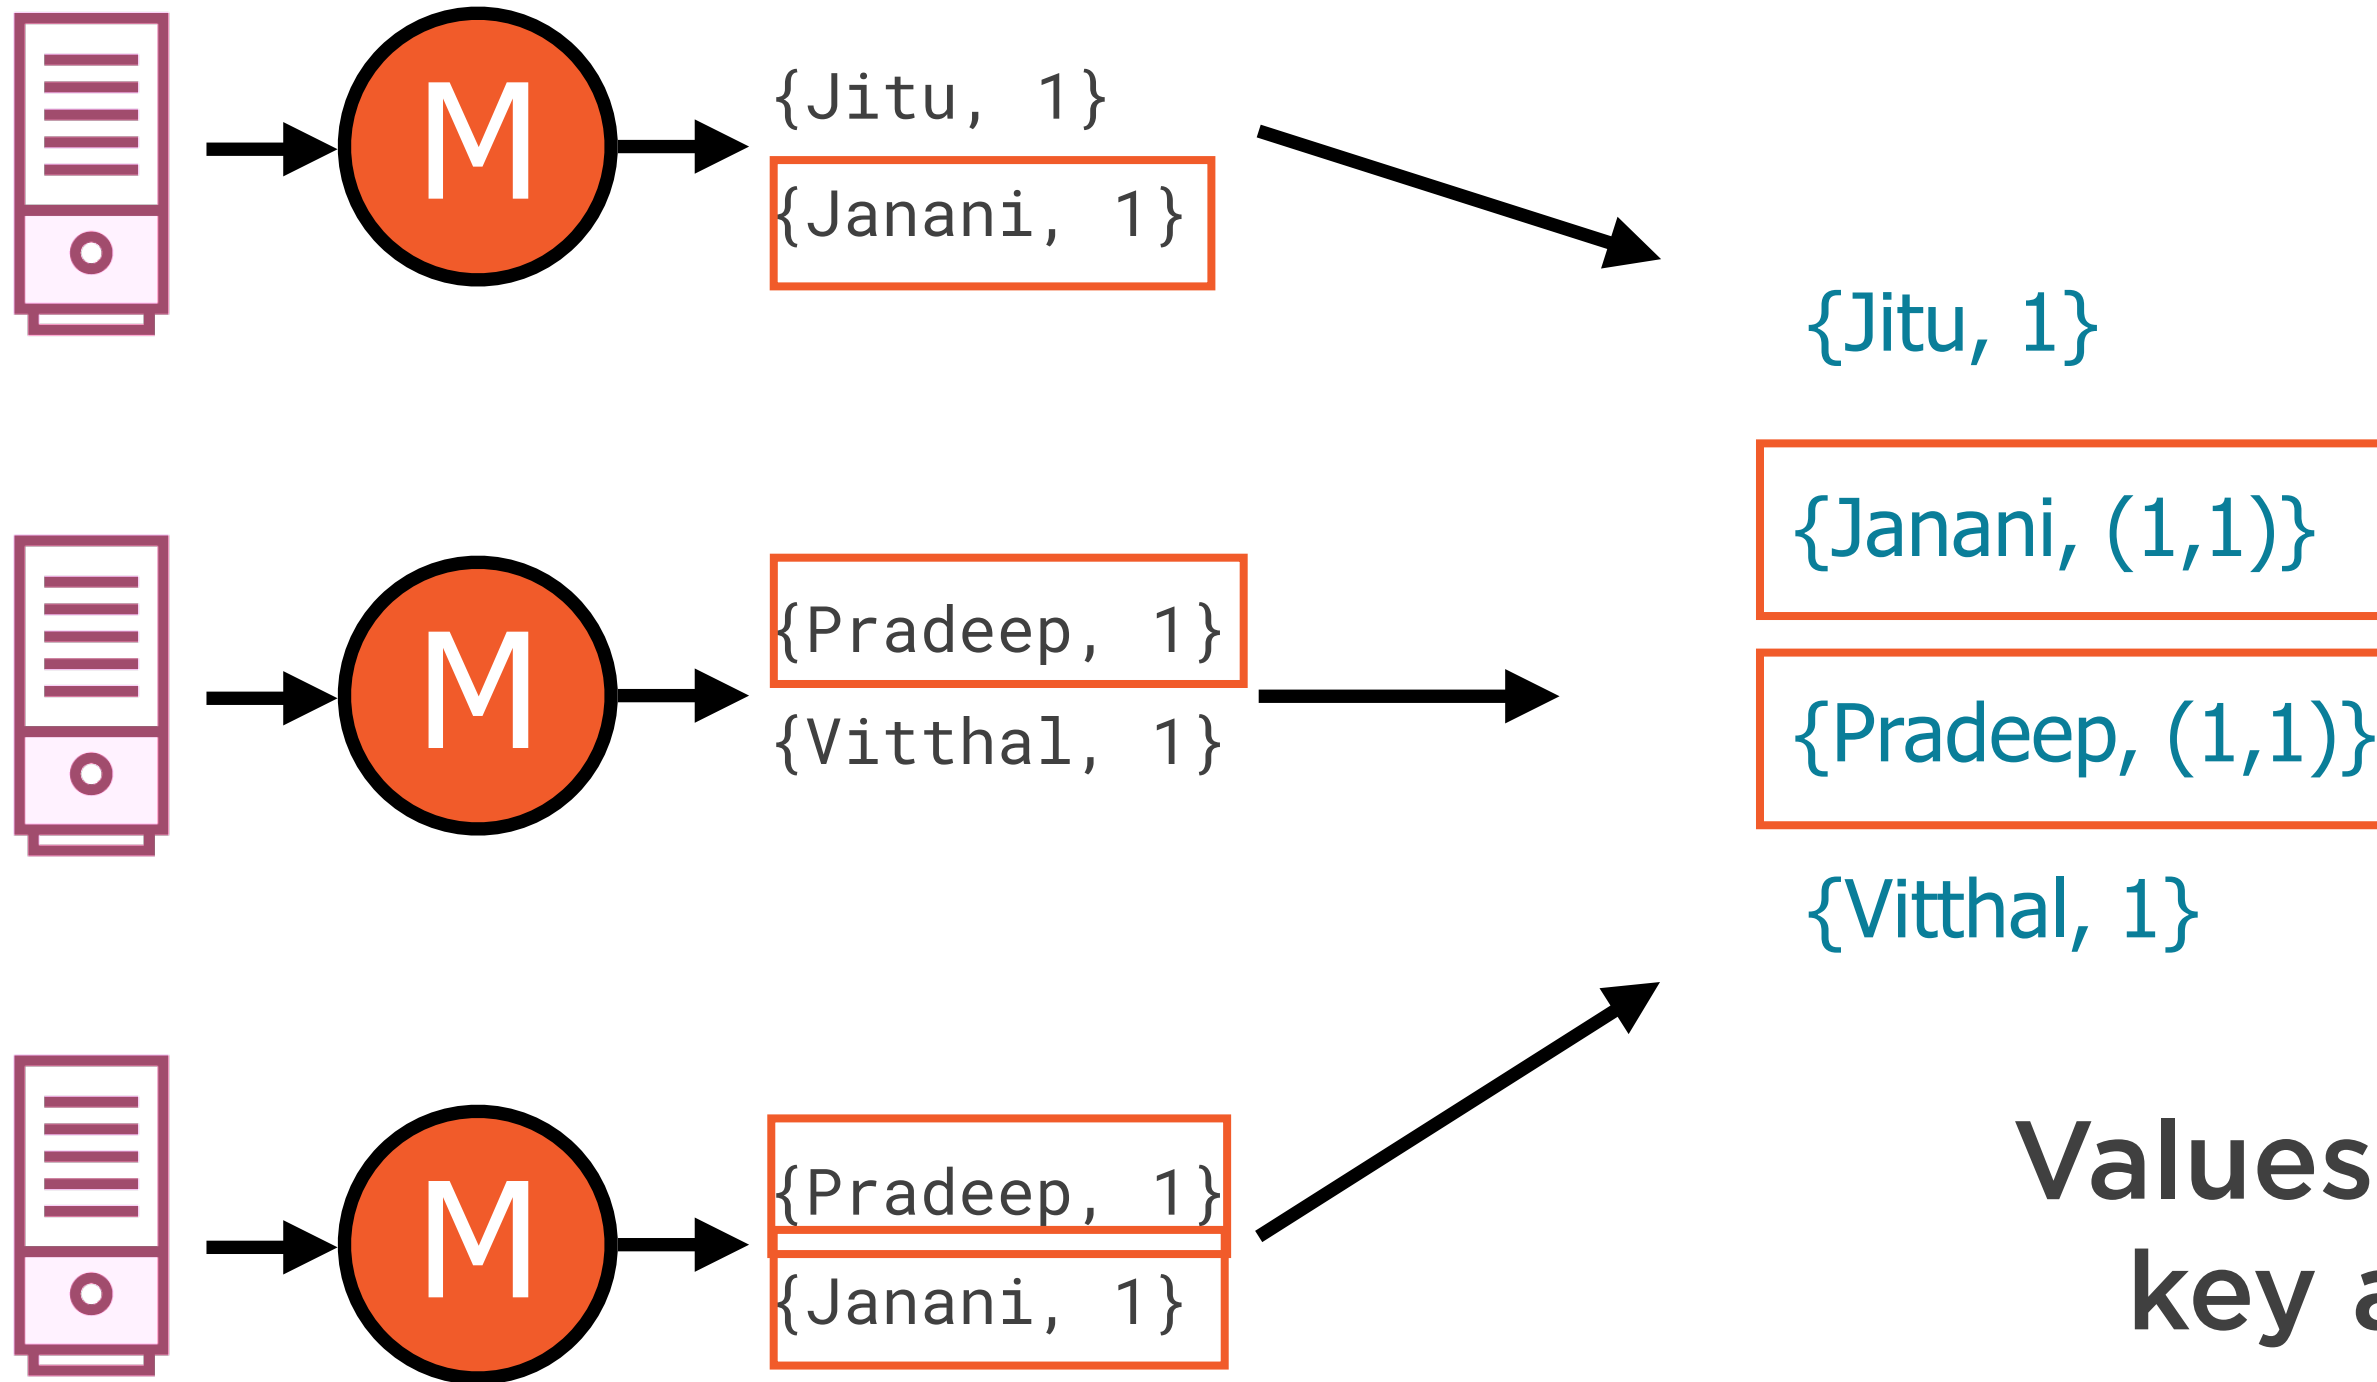
{Pradeep, 1}
{Janani, 1}

Reduce Flow



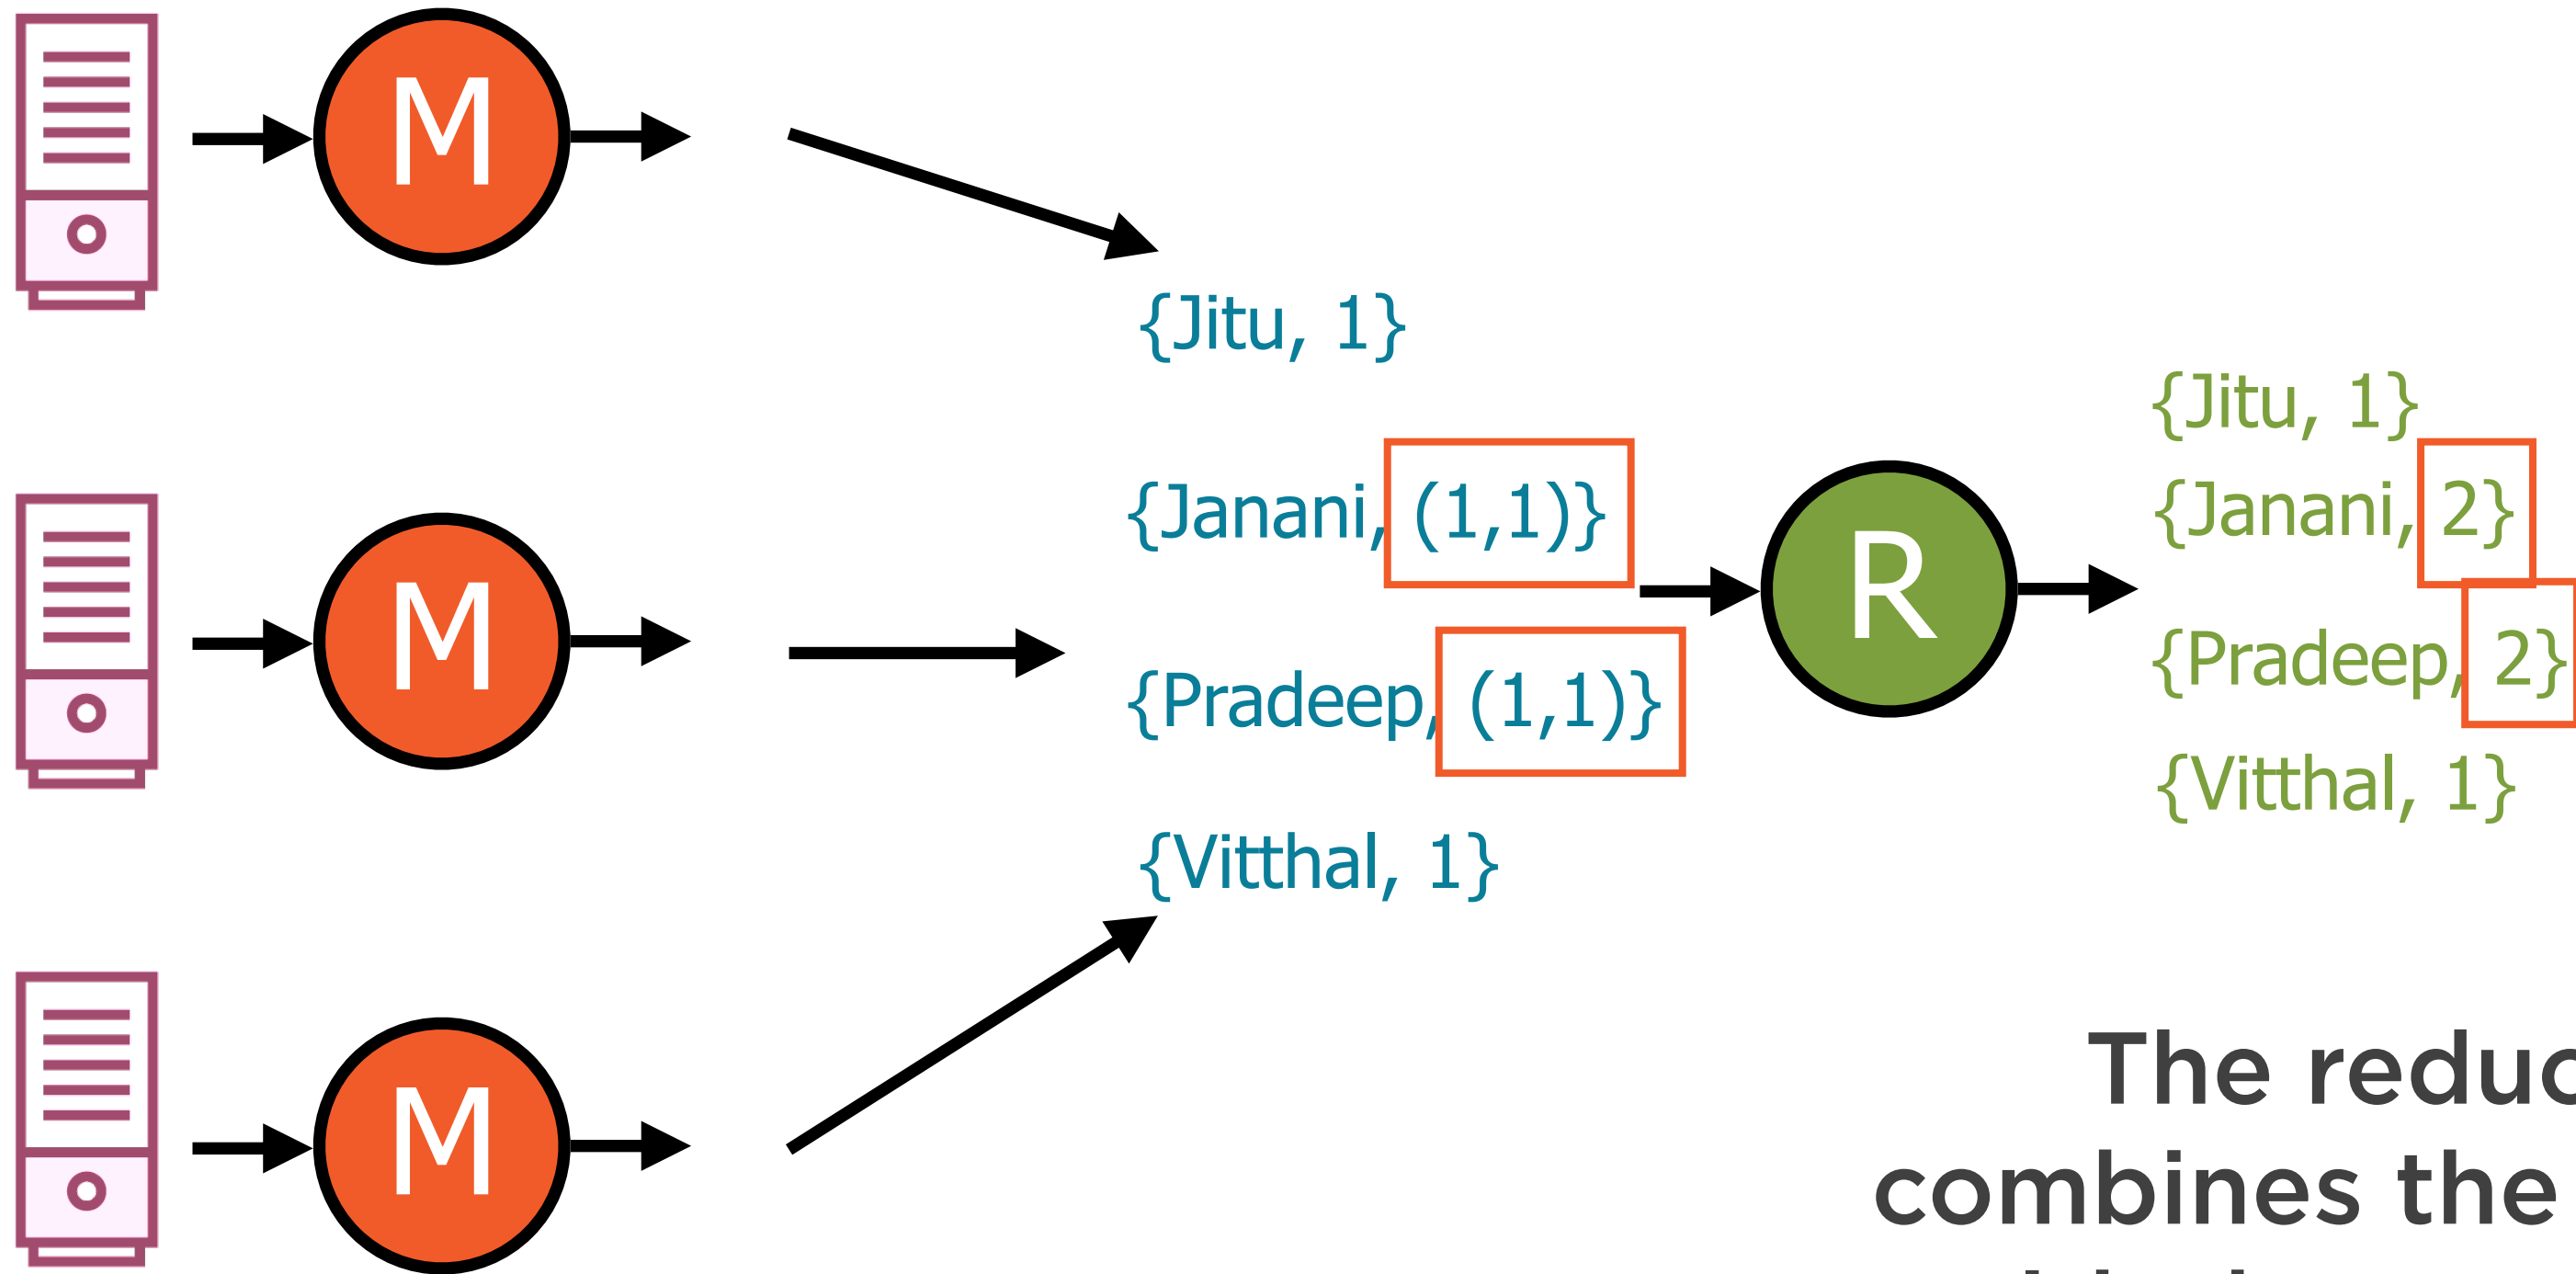
The results are
passed on to another
process ie. a **reducer**

Reduce Flow



**Values with the same
key are collected
together**


Reduce Flow



**The reducer
combines the values
with the same key**

MapReduce Flow

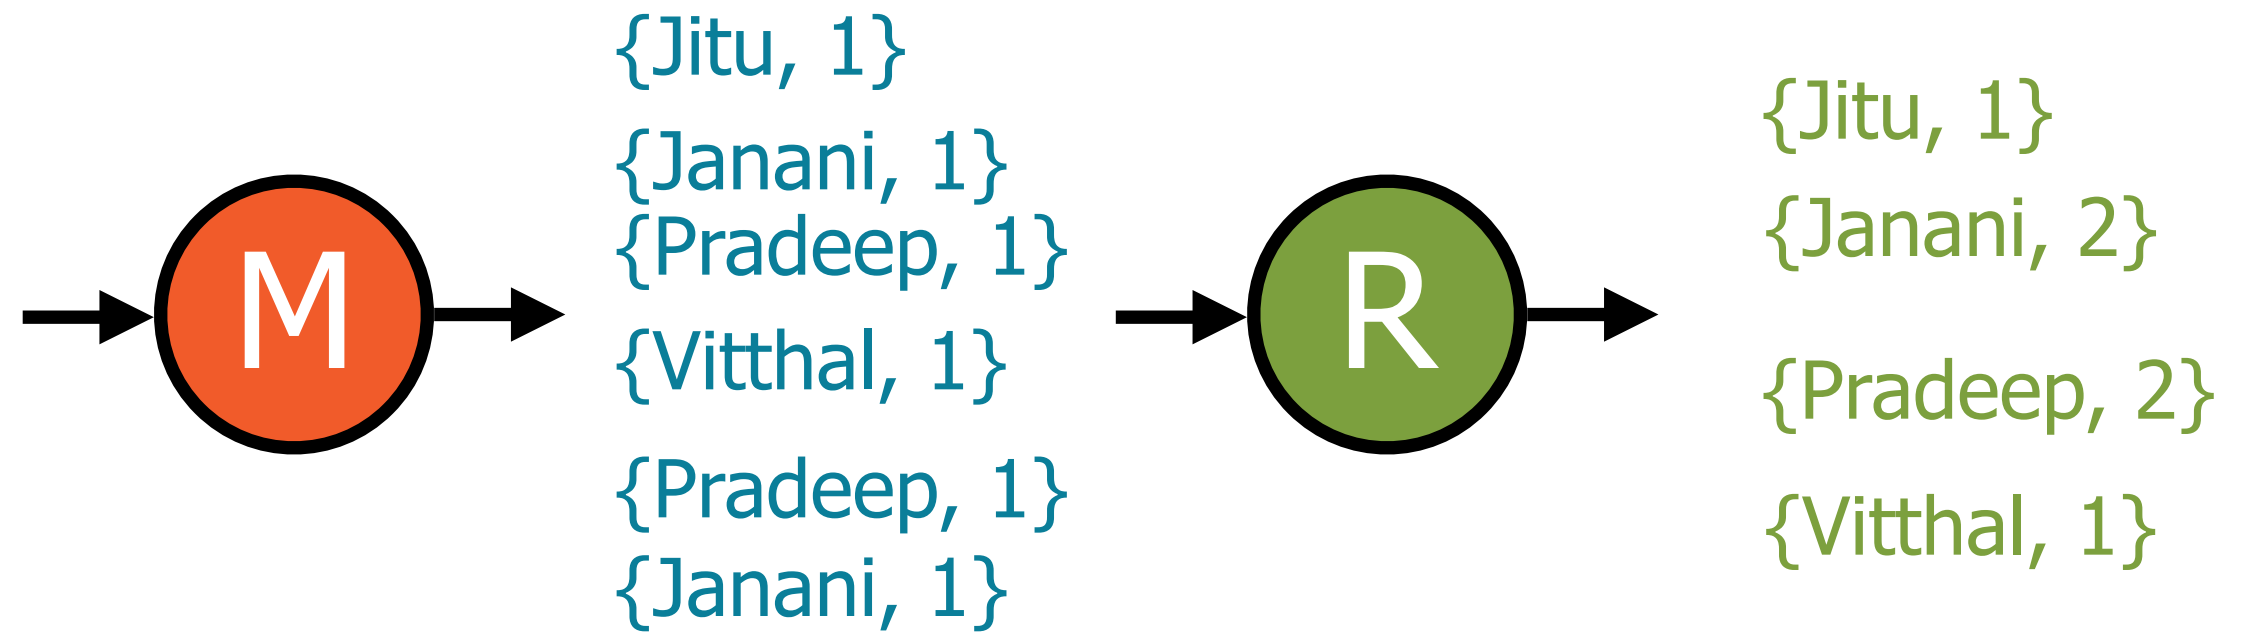
View	From	To
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep



The basic form of
EVERY MapReduce task

MapReduce Flow

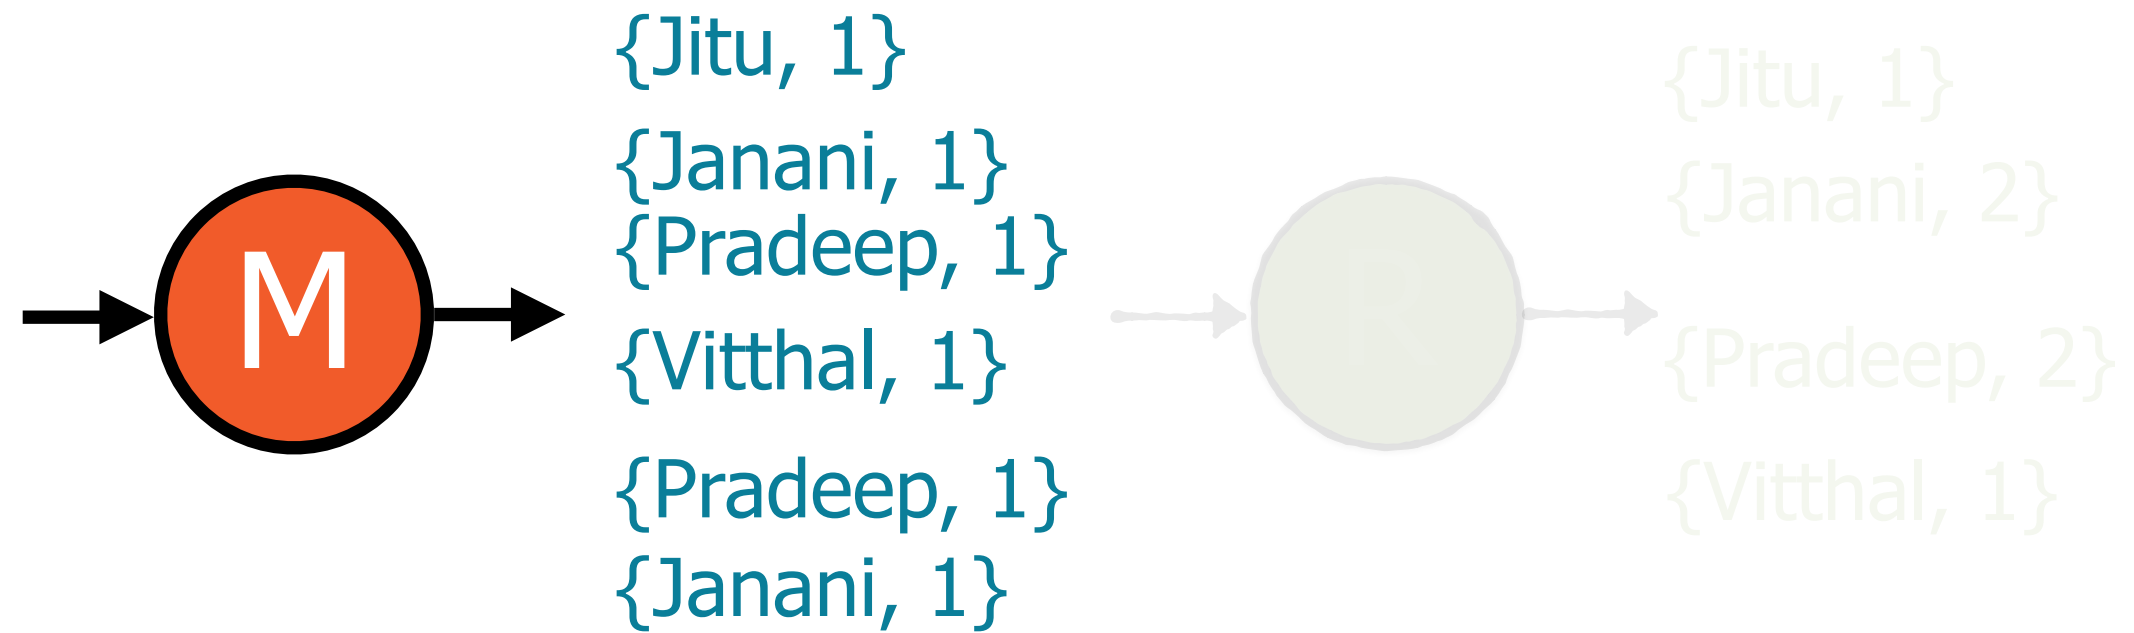
View	From	To
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep



The basic form of
EVERY MapReduce task

MapReduce Flow

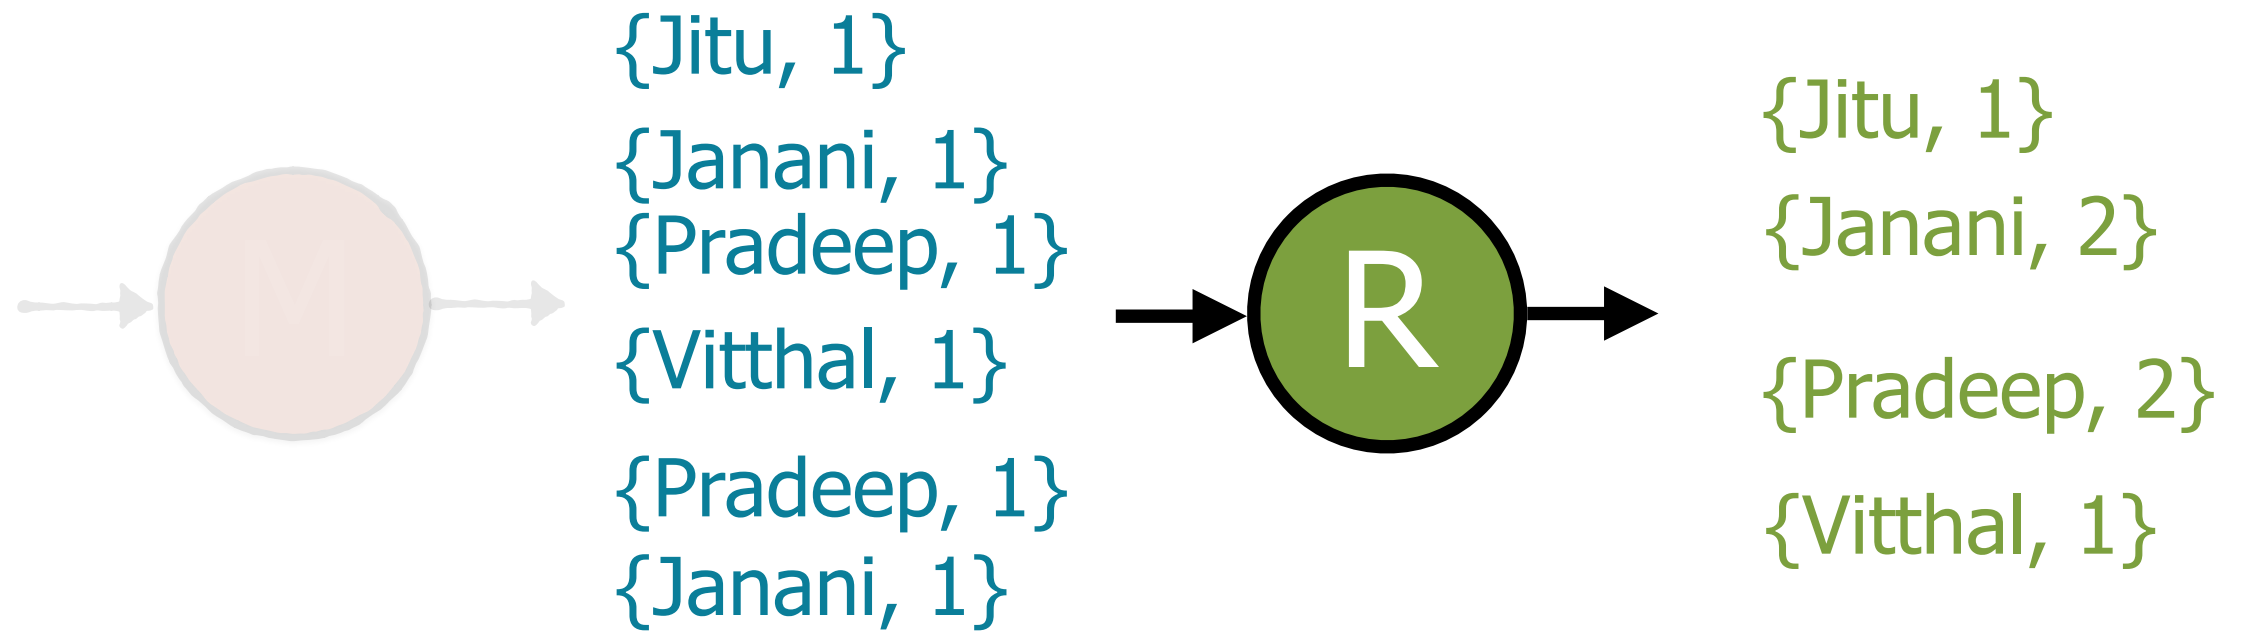
View	From	To
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep



The Map step produces a set of {key, value} pairs

MapReduce Flow

View	From	To
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep




The Reduce step combines {key, value} pairs with the same key

Key Insight Behind MapReduce



Many data processing tasks can be expressed in this form

Answer 2 Questions




1. What {key, value} pairs should be emitted in the map step?
2. How should values with the same key be combined?

Building a User-ViewCount Map

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep

**How many views has
each profile had?**

Building a User-ViewCount Map




Filter Views for a User

View ID	From Member	To Member
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep

**What information regarding
page views should be
shown to one user?**

Filter Views for a User

View	From	To
1	Janani	Jitu
2	Swetha	Janani
3	Shreya	Pradeep
4	Jitu	Vitthal
5	Shreya	Janani
6	Jitu	Pradeep



View ID	From	To
3	Shreya	Pradeep
6	Jitu	Pradeep

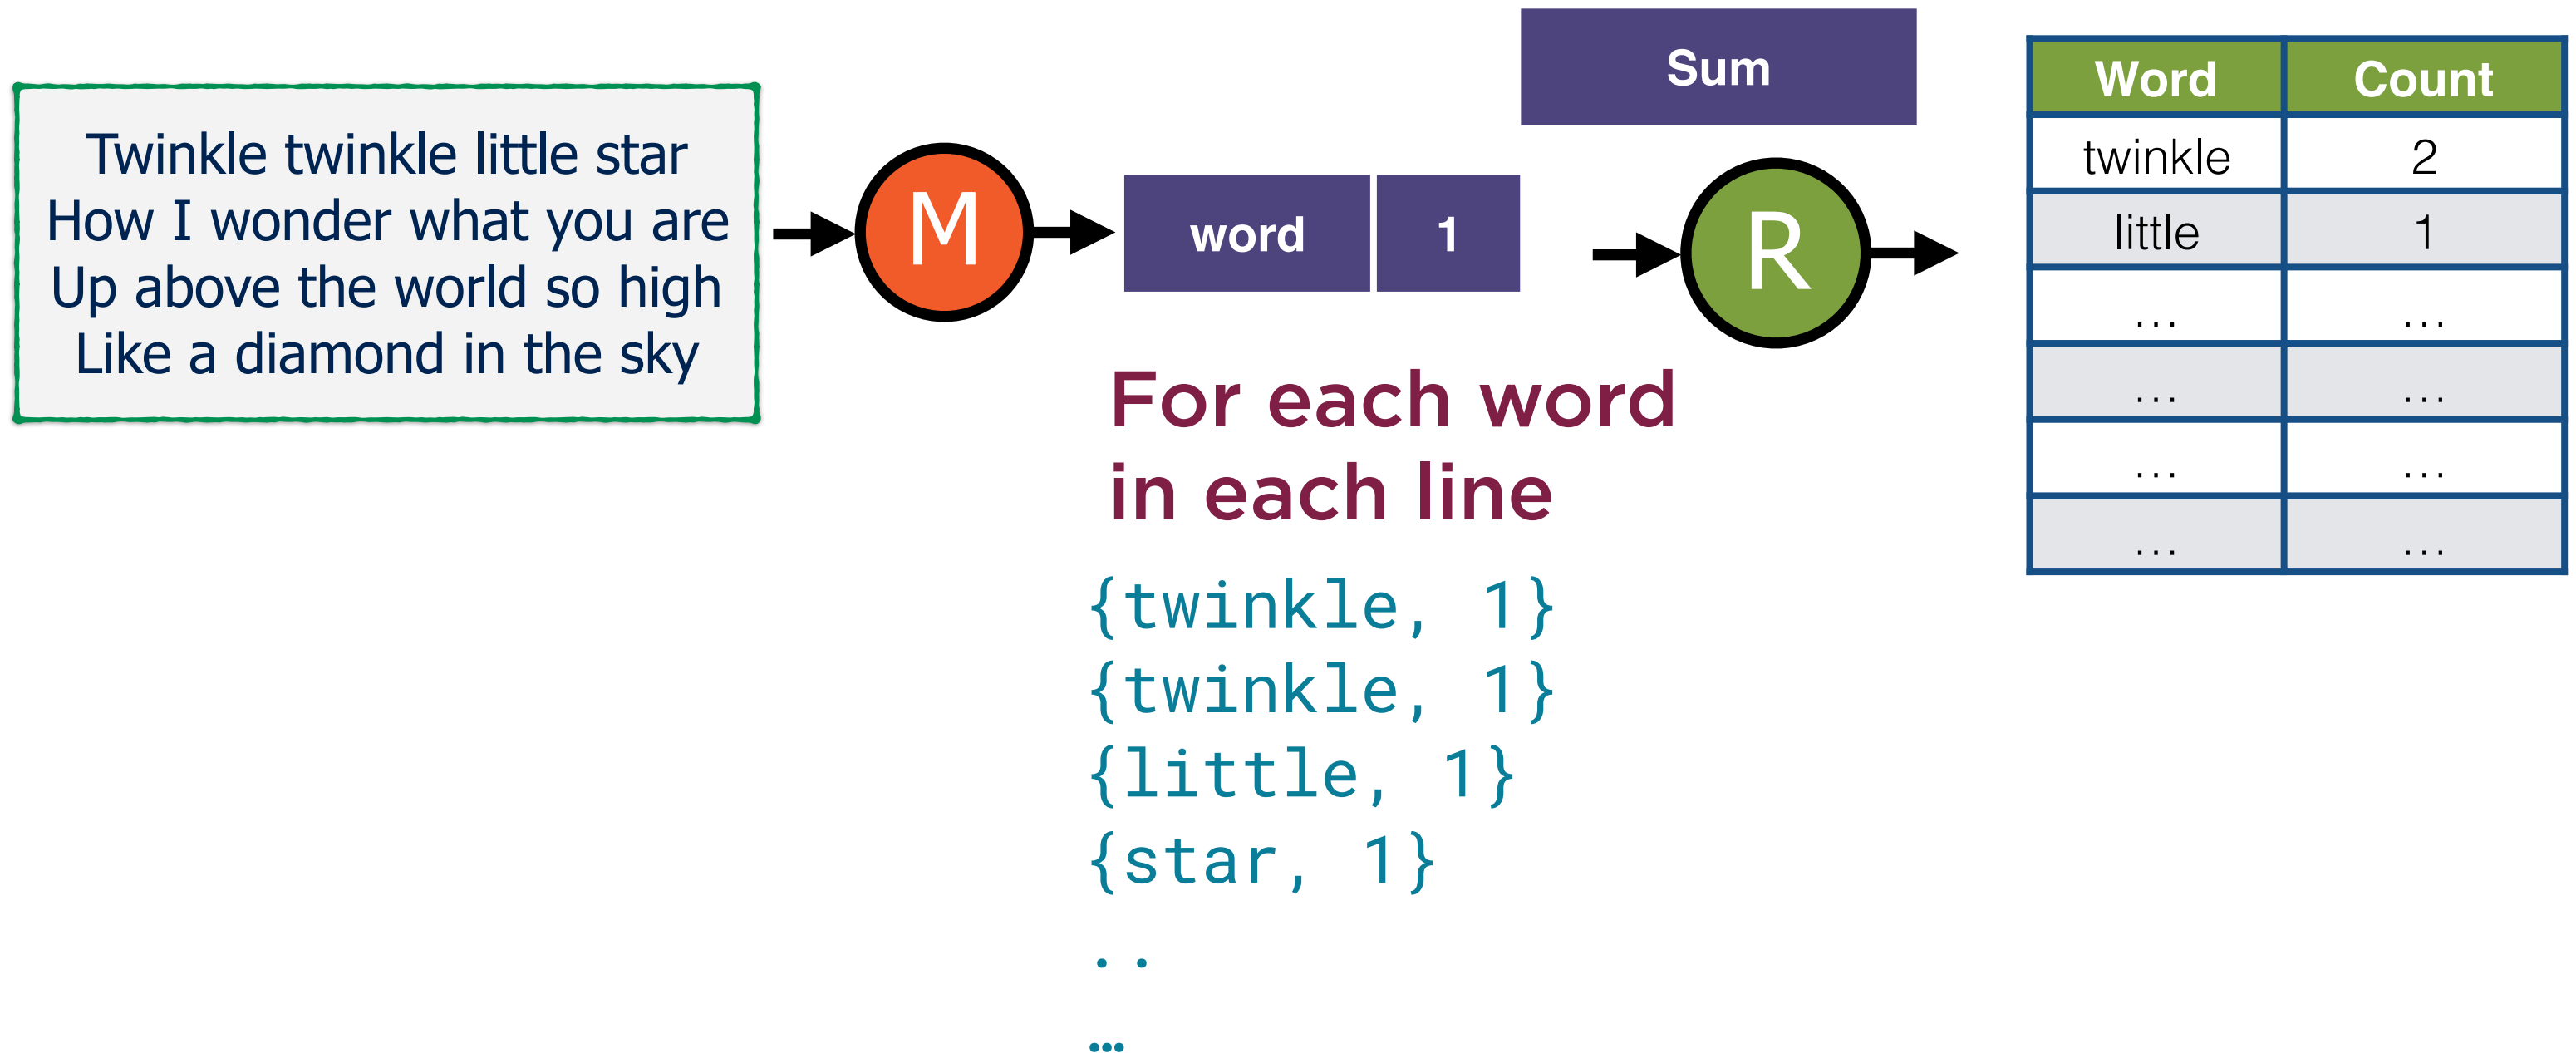
If the user
matches the
given user

Word Counts in a Document

Twinkle twinkle little star
How I wonder what you are
Up above the world so high
Like a diamond in the sky

**What is the frequency of
every word in this
document?**

Word Counts in a Document

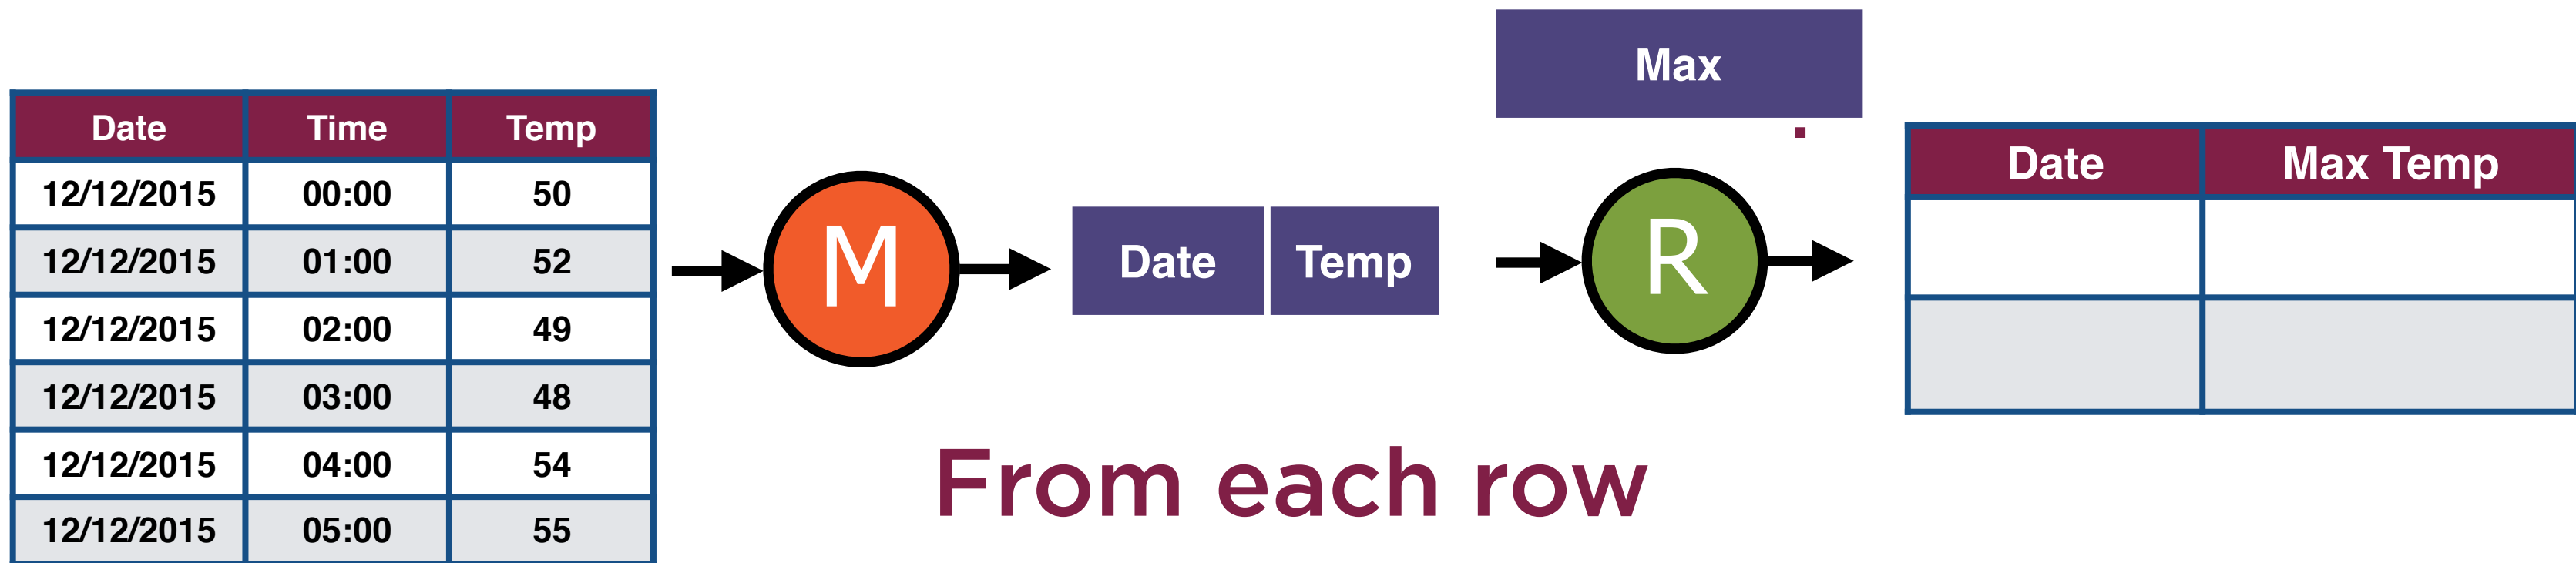


Max Temperature in a Day

Date	Time	Temp
12/12/2015	00:00	50
12/12/2015	01:00	52
12/12/2015	02:00	49
12/12/2015	03:00	48
12/12/2015	04:00	54
12/12/2015	05:00	55

What is the maximum
recorded temperature on a
given day?

Max Temperature in a Day





Answer these to
parallelize any task :)

Summary

Understand the need for Distributed Computing

Understand the role of MapReduce in a distributed computing setup

Spot applications of MapReduce

Know the typical flow of a MapReduce task